

Important Configuration



<https://www.elastic.co/guide/en/elasticsearch/reference/current/important-settings.html>



1. Path of Data/Log



2. Cluster name



3. Node name



4. bootstrap.memory_lock: true

Try to lock the process address space into RAM



5. network.host
not use loopback address



6. discovery.zen.ping.unicast.hosts



7. discovery.zen.minimum_master_nodes

to prevent data loss

$(\text{master node} / 2) + 1$



8. DELETE ALL !!
please disable



9. Increase index.refresh_interval

default = 1s



10. Disable index.refresh_interval

set value to -1

Load a large amount of data at once



11. Index buffer size

default = 10%



12. Use bulk request



13. Use multiple worker/thread to send data to Elasticsearch



Use faster hardware !!



Bootstrap Checks



<https://www.elastic.co/guide/en/elasticsearch/reference/current/bootstrap-checks.html>



Heap size 50%

File descriptor => ulimit

Maximum number of thread => +2048

Disable swapping

Virtual Memory => vm.max_map_count=262144

